Data2Vis: Automatic Generation of Data Visualizations Using Sequence to Sequence Recurrent Neural Networks

Victor Dibia [*]
IBM Research

Çağatay Demiralp [†]
MIT CSAIL & Fitnescity Labs

ABSTRACT

Rapidly creating effective visualizations using expressive grammars is challenging for users who have limited time and limited skills in statistics and data visualization. Even high-level, dedicated visualization tools often require users to manually select among data attributes, decide which transformations to apply, and specify mappings between visual encoding variables and raw or transformed attributes. In this paper we introduce Data2Vis, an end-to-end trainable neural translation model for automatically generating visualizations from given datasets. We formulate visualization generation as a language translation problem where data specifications are mapped to visualization specifications in a declarative language (Vega-Lite). To this end, we train a multilayered attention-based encoder-decoder network with long short-term memory (LSTM) units on a corpus of visualization specifications. Qualitative results show that our model learns the vocabulary and syntax for a valid visualization specification, appropriate transformations (count, bins, mean) and how to use common data selection patterns that occur within data visualizations. Data2Vis generates visualizations that are comparable to manually-created visualizations in a fraction of the time, with potential to learn more complex visualization strategies at scale.

Index Terms: Human-centered computing—Visualization—Visualization techniques—Treemaps; Human-centered computing—Visualization—Visualization design and evaluation methods

1 INTRODUCTION

Users create data visualizations using a range of tools with a range of characteristics (Figure 1). Some of these tools are more expressive, giving expert users more control, while others are easier to learn and faster to create visualizations, appealing to general audiences. For instance, imperative APIs such as OpenGL and HTML Canvas provide greater expressivity and flexibility but require significant programming skills and effort. On the other hand, dedicated visual analysis tools and spreadsheet applications (e.g., Microsoft Excel, Google Spreadsheets) provide ease of use and speed in creating standard charts based on templates but offer limited expressivity and customization.

Declarative specification grammars such as ggplot2 [71], D3 [10], Vega [58], and Vega-Lite [57] provide a trade-off between speed and expressivity. However, these grammars also come with steep learning curves, can be tedious to specify depending on the syntax and abstraction level adopted, and can suffer from reusability issues. In fact, there is little known about the developer experience with visualization grammars, beyond the degree with which they are used. For example, ggplot2 can be difficult for users who are not familiar with R. Vega, which is based on a JSON schema, can be tedious even for users who are familiar with JSON. Even tools with higher-level abstractions such as the ones based on chart templates often require

---

[*]e-mail: dibiavc@us.ibm.com

[†]e-mail: cagatay@csail.mit.edu

Figure 1: Axis of visualization specification. Data visualizations are created with a spectrum of tools with a spectrum of speed and expressivity. Some of these tools are faster but others are more expressive to create visualizations.

the user to manually select among data attributes, decide which statistical computations to apply, and specify mappings between visual encoding variables and either the raw data or the computational summaries. This task can be daunting with complex datasets especially for typical users who have limited time and limited skills in statistics and data visualization. To address these challenges, researchers have proposed techniques and tools to automate designing effective visualizations [14, 19, 39, 40, 47, 54] and guide users in visual data exploration [2, 18, 25, 48, 54, 61, 63, 69, 75, 77, 78].

Prior techniques and tools for automated visualization design and visualization recommendation are based on rules and heuristics. The need to explicitly enumerate rules or heuristics limits the application scalability of these approaches and does not take advantage of expertise codified within existing visualizations. Automated and guided visualization design and exploration can significantly benefit from implicitly learning these rules from examples (i.e., data), effectively incorporating both data and visualization design context.

In this work, we formulate visualization design as a problem of translation between data specification and visualization specification. To operationalize our formulation, we train an LSTM-based neural translation model (Data2Vis) on a corpus [52] of Vega-Lite visualization specifications, taking advantage of Vega-Lite's (and of similar grammars') design motivation to support programmatic generation. We demonstrate the model's use in automatically generating visualizations with applications in easing the visualization authoring process for novice users and helping more experienced users jump start visualization design. Our contributions include 1) formulating visualization design as a sequence to sequence translation problem, 2) demonstrating its viability by training a sequence to sequence model, Data2Vis, on a relatively small training dataset and then effectively generating visualizations of test data, and 3) integrating Data2Vis into a web-based application that has been made publicly available at `http://hci.stanford.edu/~cagatay/data2vis`. Our work is the first in applying deep neural translation to visualization generation and has important implications for future work, opening the way to implicitly learn visualization design and visual analysis rules from examples at scale.

In what follows, we first summarize related work followed by details of the Data2Vis model and its training process. We then present our results, providing several visualization examples automatically generated using the trained model. Next we discuss the potential impact of Data2Vis and its current limitations and provide an agenda for future work. We conclude by summarizing our contributions and insights.

```
{
    "data": {
        "url": "data/stocks.csv"},
    "mark": "line",
    "encoding": {
        "x": {
            "field": "date",
            "type": "temporal",
            "axis": {
                "format": "%Y" }
        },
        "y": {
            "field": "price",
            "type": "quantitative" },
        "color": {
            "field": "symbol",
            "type": "nominal" }
    }
}
```

Figure 2: A Vega-Lite specification (left) and the generated visualization (right). Users can succinctly specify selections, transformations and interactions using the Vega-Lite grammar formatted in JSON [57].

2 RELATED WORK

Our work is related to earlier efforts in effective visualization specification, automated visualization design, and deep neural networks (DNNs) for synthesis and machine translation.

2.1 Declarative Visualization Specification

Earlier data visualization work proposes grammars and algebraic operators over data as well as visual encoding and design variables to specify visualizations (Figure 1). Wilkinson's seminal work [72] introduces a grammar of graphics and its implementation (VizML), greatly shaping the subsequent research on visualization specification. Polaris [64] (commercialized as Tableau) uses a table algebra drawn from Wilkinson's grammar of graphics. The table algebra of Polaris later evolved to VizQL [30], forming the underlying representation of Tableau visualizations. Wickham introduces ggplot2 [71], a widely-popular package in the R statistical language, based on Wilkinson's grammar. Similarly, Protovis [9], D3 [10], Vega [58], Brunel [74], and Vega-Lite [57] all provide grammars to declaratively specify visualizations. Some of them require more complete specifications than others. For instance, Protovis, D3 and Vega support finer control over visualization specification with incurred cost of verbosity.

Wongsuphasawat et al. [77] introduce Vega-Lite (Figure 2) to support Voyager, a faceted browser for visualization recommendations. Vega-Lite is a high-level grammar built on top of Vega to facilitate clarity and conciseness with some loss in expressivity. The expressivity of Vega-Lite is a strict subset of Vega. We train our model on a Vega-Lite corpus [52], which contains datasets and corresponding visualizations specified in Vega-Lite.

Declarative grammars eschew chart templates typically used in dedicated visualization tools or spreadsheet applications such as Microsoft Excel and Google Spreadsheets, which have limited support for customization. Conversely, these grammars facilitate expressivity by enabling a combinatorial composition of low-level building blocks such as graphical marks, scales, visual encoding variables, and guides. However, increased expressivity often decreases the speed with which visualizations can be created and makes the learning more difficult, limiting the number of users who can effectively use the specification method. One of our aims with Data2Vis is to bridge this gap between the speed and expressivity in specifying visualizations.

2.2 Automated Visualization

Prior work proposes desiderata and tools (e.g., [14, 19, 40, 47, 54]) to automatically design effective visualizations, building on Bertin's study [7] of visual encoding variables and earlier graphical perception research, e.g., [1,5,17,41,47,62]. Earlier research also develops interactive systems and recommendation schemes [11, 27, 48, 61, 63, 67–69, 73, 75, 77, 78] to guide users in exploratory data analysis and visualization design. PRIM-9 [25], GrandTour [2] SeeDB [69], Zenvisage [63], ShowMe [48], Voyager [77], Voyager 2 [78], SAGE [54] and VizDeck [39] prioritize charts according to one or more evaluation measures such as data saliency, data coverage, perceptual effectiveness, user task, and user preferences. Similarly, Rank-by-Feature [61], AutoVis [75], and Foresight [18] use statistical criteria over data attributes and instances in recommending and ranking visualizations.

Data2Vis represents a departure from rule-based approaches of prior work both in conceptual formulation and technical approach taken. It makes contributions by specifying how automated visualization can be cast as a learning problem, providing a concrete implementation of a deep learning model for visualization generation. Data2Vis emphasizes the creation of visualizations specifications using rules *learned* from examples, without resorting to a predefined enumeration of rules or heuristics, complementing earlier work. Researchers recently recognized the potential of machine learning in automating visualization design and visual analysis [56], applying machine learning for recommending visualizations [32, 43, 55] and refining visualization recommendations [49]. Data2Vis differs from this exciting line of recent work, which relies on feature extraction and manual constraint specification, in learning to automatically generate visualizations from data with an end-to-end approach.

Adopting a learning approach to designing automated visualization systems holds potential for improving the maintenance and scalability of such systems. Existing approaches are limited by a dependence on a set of manually created (interdependent) rules which can be voluminous, tedious update, and may not sufficiently cover edge cases necessary to generate good visualizations. By using a learning approach, we avoid these limitations as a learned model can better represent the visualization rule space given sufficient examples. Further more, the performance and capabilities of the system can be improved by improving the dataset of examples used to train models within learning based systems. As more users author visualizations, the system can leverage *experiences and rules* encoded within these visualizations, to increase it's coverage and *scale* its performance. The visualization generation capabilities of Data2Vis can also be integrated into existing higher-level recommendation systems of visual data exploration and used in tandem with rule-based techniques to drive these systems. We published the current work earlier as a preprint [22] and made the source code for the Data2Vis model publicly available [21].

2.3 Deep Neural Networks for Synthesis

Prior deep neural network (DNN) research studies adopt generative approaches to learn human-like cognitive and creative capabilities. Examples include the use of models to synthesize music, drawings, images from textual descriptions, code from hand-drawn sketches or interface screenshots. Ha et al. [29] train a recurrent neural network (RNN) to predict and generate stroke-based drawings of common objects. Reed et al. [53] present a DNN architecture and generative adversarial network (GAN) formulation to "translate" textual visual concepts to pixels. Others learn how to generate code from user interface screenshots [6] and how to compose music using purely sequential models [24, 34] and cascading a sequential model with a restricted Boltzman machine [12]. All these approaches aim to simplify the creative process for both novices and experts. In this sense, our work here shares a motivation with prior work. We also use a variation of sequential neural network models, a sequence to sequence model, to generate visualization specifications from given data.
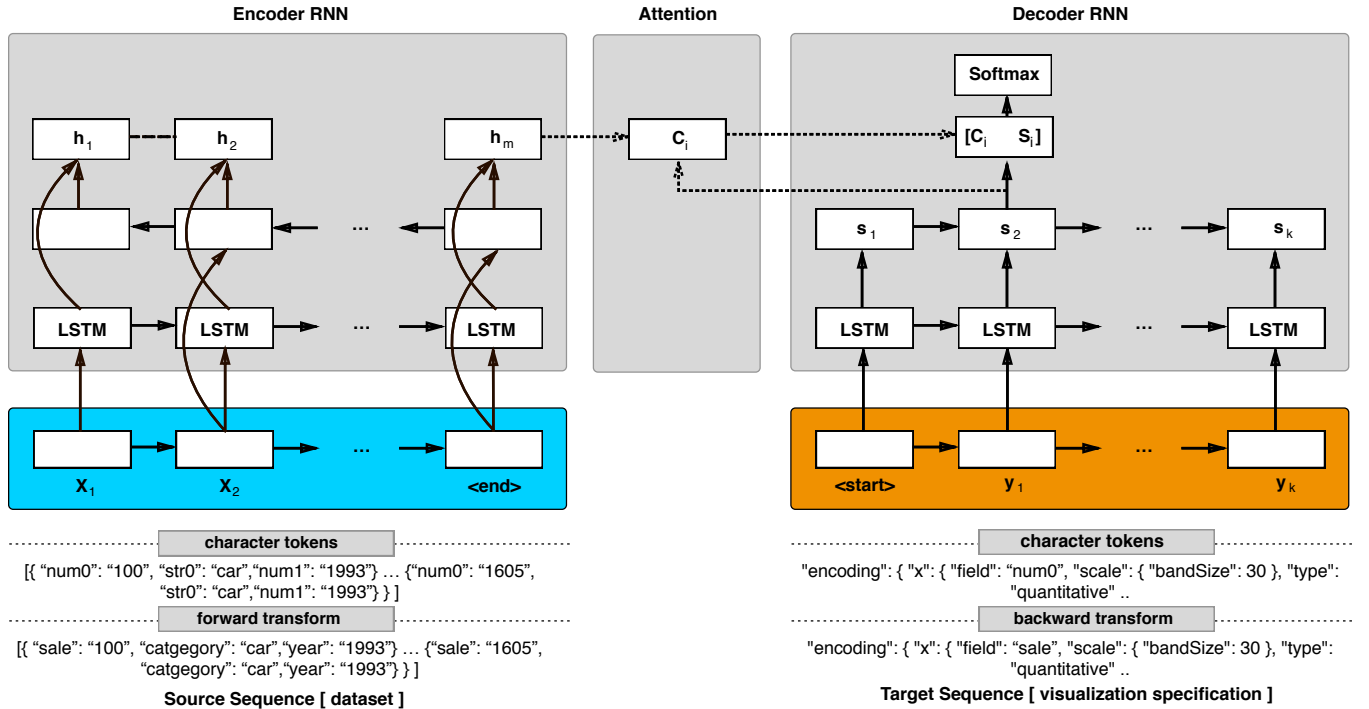
Figure 3: Data2Vis is a sequence to sequence model with encoder-decoder architecture and attention module. To simplify learning, we perform simple forward and backward transformations on the source (dataset in JSON format) and target sequence (Vega-Lite visualization specification) which are then converted to character tokens.

2.4 Deep Neural Networks for Machine Translation

Recent work introduces DNN models, e.g., [3,16,35,45,66] that significantly improves [33,46,60,79] the performance of machine translation systems, surpassing the preceding phrase-based approaches. Deep neural translation models eschew hand engineering the features, in large part, by using large training data, enabling the end-to-end learning in practice. Sequence to sequence models (e.g., [3,45]) are a particularly successful and popular class of deep learning models applied in machine translation (see [13] for an evaluation of alternative architectures). Akin to autoencoders, these models have also a symmetric, encoder-decoder architecture. Sequence to sequence models are composed of encoder-decoder layers which consist of recurrent neural networks (RNNs) and an attention mechanism that aligns target tokens with source tokens.

In addition to translating between natural languages, earlier work, e.g., [4, 15, 20, 42, 51, 81] also uses DNNs to translate between two domain specific languages (DSLs), between a natural language specification and a DSL (e.g. translating from natural language to SQL [23,82]), and between two programming languages. Similar to the prior work translating between general or domain specific programming languages, Data2Vis also translates between two formal languages. Ling et al. [42] use a sequence to sequence model to translate TCG (Trading Card Games) cards to their Python and Java specifications without explicitly representing the target syntax. Data2Vis is also a sequence to sequence model that directly uses textual source and target specifications without representing their syntax (e.g., using abstract syntax trees) explicitly.

3 PROBLEM FORMULATION

Building on earlier work that applies deep learning for translation and synthesis, we formulate the data visualization problem as a sequence to sequence translation problem, which can be readily addressed using sequence to sequence models (seq2seq) [3,16,66]. Our input sequence is a dataset (fields, values in json format) and


Figure 4: Frequency of the Vega-Lite mark types and transforms used in our training examples.

our output sequence is a valid Vega-Lite [57,58] visualization specification.

Existing models used for sequence translation [3, 13, 16, 44, 45, 66] belong to a family of *encoder-decoder* networks where the *encoder* reads and encodes a source sequence into a fixed length vector, and a *decoder* outputs a translation based on this vector. The entire encoder-decoder system is then jointly trained to maximize the probability of outputting a correct translation, given a source sequence.

While sequence to sequence models have originally focused on generating data that is sequential or temporally dependent e.g language translation [3, 13, 16, 44, 45, 66], they also find applications for problems where the output or input is non-sequential as seen in text summarization [50, 65] and image captioning [36, 70, 80]. Two important advances that enable non-sequential use cases include the introduction of bidirectional RNN units [59] and attention mechanisms [3, 70, 80]. An ordinary RNN (unidirectional) reads an input sequence x from the frist token x_1 to the last x_m and generates an encoding only based on the preceding tokens it has seen. On

the other hand, a Bidirectional RNN (BiRNN) consists of both a forward RNN and a backward RNN, which enables an encoding generation based on both the preceding and following tokens. The forward RNN $\overrightarrow{f}$ reads the input sequence as it is ordered (from x_1 to x_m) and calculates a sequence of forward hidden states ($\overrightarrow{h_1}$, ... ,$\overrightarrow{h_m}$). The backward RNN $\overleftarrow{f}$ reads the sequence in the reverse order (from x_m to x_1), resulting in a sequence of backward hidden states ($\overleftarrow{h_1}$, ... ,$\overleftarrow{h_m}$). Thus, when a BiRNN is used to encode an input sequence, it generates a hidden state $\overrightarrow{h_j}$ which is a concatenation of both the forward and backward RNNs, $h_j = \left[\overrightarrow{h_j}^{\top} ; \overleftarrow{h_j}^{\top} \right]^{\top}$ and contains summaries of both the preceeding and following tokens. Attention mechanisms allow a model to focus on aspects of an input sequence while generating output tokens. They provide the additional benefits of making translation models robust to performance degradation while generating lengthy sequences, and enable the model to learn mappings between source and target sequences of different lengths [3]. For example, when used in image captioning, attention mechanisms allow the model to focus on specific parts of objects in an image, while generating each word or token in the image caption. Furthermore, attention mechanisms improve our ability to interpret and debug sequence to sequence models as they provide valuable insights on *why* a given token is generated at each step. Taken together, these two important advances enable us to use a sequence translation model that first takes into consideration the entire data input (dataset) and then focus on aspects of the input (fields) in generating a visualization specification.

Seq2seq models for language translation are trained using embeddings of the source and target tokens which can be generated based on words, subword or per character units [3, 16, 66]. We select a per character unit tokenization given our source and target sequences consist of symbols as opposed to learnable word groups seen in related problems like language translation.

4 MODEL

Our model (Figure 3) is based on an encoder-decoder architecture with attention mechanism that has been previously applied in machine translation [3, 44, 45]. The encoder is a bidirectional recurrent neural network (RNN) that takes in an input sequence of source tokens $x = (x_1, ..., x_m)$ and outputs a sequence of states $h = (h_1, ..., h_m)$. The decoder is also an RNN that computes the probability of a target sequence $y = (y_1, ..., y_k)$ based on the hidden state h. The probability of each token in the target sequence is generated based on the recurrent state of the decoder RNN, previous tokens in the target sequence and a context vector c_i. The context vector (also called the attention vector) is a weighted average of the source states and designed to capture the context of source sequence that help predict the current target token.

We use a 2-layer bidirectional RNN encoder and a 2-layer RNN decoder, each with 512 Long Short-Term Memory (LSTM) [26, 31] units (cells). To decide which RNN unit type to use, we experimented with both gated recurrent unit (GRU) [16] and LSTM, both of which are common RNN cell variants. We found LSTM cells provided better results (valid json, valid Vega-Lite specification) compared to GRU cells, which concurs with earlier empirical results [13].

5 DATA AND PREPROCESSING

To generate plausible visualizations conditioned on a given source dataset, our model should achieve several learning objectives. First, the model must select a subset of fields to focus on when creating visualizations (most datasets have multiple fields which cannot all be simultaneously visualized). Next, the model must learn differences in data types across the data fields (numeric, string, temporal, ordinal, categorical etc.), which in turn guides how each field is specified in the generation of a visualization specification. Finally, the model

must learn the appropriate transformations to apply to a field given its data type (e.g., aggregate transform does not apply to string fields). In our case, this includes view-level transforms (aggregate, bin, calculate, filter, timeUnit) and field level transforms (aggregate, bin, sort, timeUnit) supported by the Vega-Lite grammar.

Achieving these objectives using a character based sequence model can be resource intensive. While character based models result in smaller vocabulary size and are more accurate for specialized domains, they also present challenges — a character tokenization strategy requires more units to represent a sequence and requires a large amount of hidden layers as well as parameters to model long term dependencies [8]. To address this issue and scaffold the learning process, we perform a set of transformations. First, we replace string and numeric field names using a short notation — "str" and "num" in the source sequence (dataset). Next, a similar backward transformation (post processing) is eplicated in the target sequence to maintain consistency in field names (see Figure 3). These transformations help scaffold the learning process by reducing the vocabulary size, and prevents the LSTM from learning field names (as we observed in early experiments). In turn we are able to reduce the overall source and target sequence length, reduce training time and reduce the number of hidden layers which the model needs to converge. Our training dataset is constructed from 4300 Vega-Lite visualizations examples, based on 11 distinct datasets. The examples were originally compiled by [52] where the authors use the CompassQL [76] recommendation engine within Voyager2 [77] to generate charts with 1-3 variables, filtered to remove problematic instances. These charts are generated based on heuristics and rules which enumerate, cluster and rank visualizations according to data properties and perceptual principles [77]. While these examples contain a simplified range of transformations and do not encode any interactions, they represent valid Vega-Lite examples and conform to important perceptual principles enforced by rules within Voyager2. These characteristics make the dataset a suitable, low-complexity test bed for benchmarking our model's performance on the task of learning to generate visualizations given only input data.

Similar to datasets observed in the wild, our sample dataset contains charts with 6 different types of visualizations (area, bar, circle, line, point, tick) and three different transforms (aggregate, bin, timeUnit)(see Figure 4). Based on this similarity, we expect similar learning performance when our model is trained with real world data sets. To generate our training dataset, we iteratively generate a source (a single row from the dataset) and target pair (see Figure 3) from each example file. Each example is then sampled 50 times (50 different data rows with the same Vega-Lite specification) resulting in a total of 215,000 pairs which are then used to train our model.

5.1 Training

We begin by generating a character vocabulary for our source and target sequences (84 and 45 symbols, respectively). A dropout rate of 0.5 is applied at the input of each cell and a maximum source and target sequence length of 500 is used. The entire model is then trained end-to-end using a fixed learning rate of 0.0001 with Adam optimizer, minimizing the negative log likelihood of the target characters using stochastic gradient descent. Our implementation is adapted from an open source neural machine translation framework by Britz *et al.* [13]. We train our model for a total of 20,000 steps, using a batch size of 32. We achieve a translation performance log perplexity metric score of 0.032, which suggests the model excels at predicting visualization specifications that are similar to specifications in our test set.

6 RESULTS

6.1 Examples of Automated Visualization Generation

Quantifying the performance of a generative model can be challenging. Following existing literature [34, 37, 38], we explore a

Figure 5: Example attention plots for a visualization generation case (a) Model learns to pay attention to field name "str" in generating the "string" field type applied to the field. (b) Model learns to pay attention to the field name "num0" and its value in specifying the "quantitative" field type applied to the field.

qualitative evaluation of the model's output. To evaluate the model, we use the Rdataset repository [1] (cleaned and converted to a valid JSON format) which was not included in our training. Figure 6 shows visualizations generated from a randomly selected dataset in the Rdataset collection. The range of valid univariate and multivariate visualizations produced suggests the model captures aspects of the visualization generation process. As training progresses, the model incrementally learns the vocabulary and syntax for valid Vega-Lite specifications, learning to use quotes, brackets, symbols and keywords. The model also appears to have learned to use the right type of variable specifications in the Vega-Lite grammar (e.g. it correctly assigns a string type for text fields and a quantitative for numeric fields). Qualitative results also suggest the use of appropriate transformations (bins, aggregate) on appropriate fields (e.g. means are performed on numeric fields). The model also learns about common data selection patterns that occur within visualizations and their combination with other variables to create bivariate plots. As experts create visualizations, it is common to group data by geography (country, state, sex), characteristics of individuals (citizenship status, marital status, sex) etc. Early results suggests that our model begins to learn these patterns and apply them in its generation of visualizations. For example, it learns to subset data using common ordinal fields such as responses (yes/no), sex (male/female) etc and plots these values against other fields (Figure 7). Finally, in all cases, the model generates a perfectly valid JSON file and valid Vega-Lite specification with some minor failure cases (Figure 6■).

6.2   Beam Search

To explore a variety of generated visualizations, a simple beam search decoding algorithm described in Wu et al. [79] is used. As opposed to outputting the most likely (highest probability) translation of an input sequence, beam search expands all possible next steps during generation and keeps the *k* most likely, where k is a user specified parameter known as the *beam width*. Unlike conventional language translation systems where beam search is applied mainly to improve translation quality by maximizing conditional probabilities of generated sequences [28], we also explore beam search as a way to generate a *diverse* set of candidate visualizations by outputting all parallel beam results. With beam search, we observe the model generates more diverse plots, exploring combinations of chart types and the use of multiple variables. Figure 6 shows results from beam search (beam width=15) where the model focuses on two fields from the dataset, generates univariate plots for these fields, subsets the plots by sex (male/female) and uses three chart types (bar, area, line).

---

[1]Rdatasets is a collection of 1147 datasets originally distributed alongside the statistical software environment R and some of its add-on packages.

6.3   Attention Plots

To further explore the efficacy of our model, and ascertain how well it learns to use aspects of the input data in generating visualizations, we examine plots of the attention weights (Figure 5) assigned to each predicted token. Results suggest that the model assigns non-monotonic weights to different input characters while generating the parts of the specification such as the fields used for each visualization axis, the data types assigned and the transformations applied to each field. As shown in Figure 5, the model places strong weights on the characters "num0" and its value "8" while generating the "quantitative" data type which it has assigned to an axis.

6.4   Comparison with a Visualization Recommender

We compare results from Data2Vis with results from Voyager 2 [78]. Note that while Voyager 2 recommends visualizations, it requires the user to select a set of data fields of interest (limited to two selections) and additional preferences. Thus, for the purpose of qualitative comparison, we present both tools with the same race progression dataset, and select two fields to view recommendations from Voyager 2. Qualitative results are presented in Figure 9 which demonstrate that Data2Vis generates a richer variety of charts. Visualizations generated by Data2Vis are not limited to specific constraints, demonstrating its viability for the task of generating a manageable set of visualizations based on data.

6.5   Web Application Integrating Data2Vis

To further evaluate the utility of our model, we developed a web application prototype interface (Figure 6) that supports the use case of an analyst exploring data similar to [77,78]. The interface supports three primary interactions; data import, visualization generation and visualization update. First, the analyst is able to import a dataset into the application. They can do this by using the "load dataset" button which loads a randomly selected dataset from the Rdataset repository or they can directly paste a JSON data array in the provided text field. Next, the analyst can select the "generate" button which submits the dataset to our model, receives a Vega-Lite specification (placed in a text field) and renders the plot. Finally, the analyst can update the generated specification by opening it in the Vega-Lite editor. We showed this early prototype to two visualization researchers and our observations suggest they were able to quickly build on the specifications generated by the model, making changes to field selections and transformations.

7   DISCUSSION

We presented the very first attempt to transform data to visualizations using a deep neural network and apply a neural machine translation (seq2seq) model in automating visualization generation. Below, we discuss the potential impact of our approach and limitations of the current Data2Vis model along with future research directions.

7.1   Impact and Use Case

Making Visualization Authoring Easier Providing users with little or no programming experience with the ability to rapidly create expressive data visualizations empowers users and brings data visualization into their personal workflow. Based on our early findings, Data2Vis is able to learn patterns in data visualizations that are can be generalized to a variety of real world datasets. For example, the use of categorical variables like gender, sex, location (state, country) and the tendency to focus on such variables can be learned from data. Thus, visualizations generated which encode such principles holds potential to make data visualization more accessible, speed-up the visualization authoring process and augment the visualization capabilities of all users.

Accelerating Data Exploration For visualization experts, it is likely that visualizations created by Data2Vis may be insufficient for
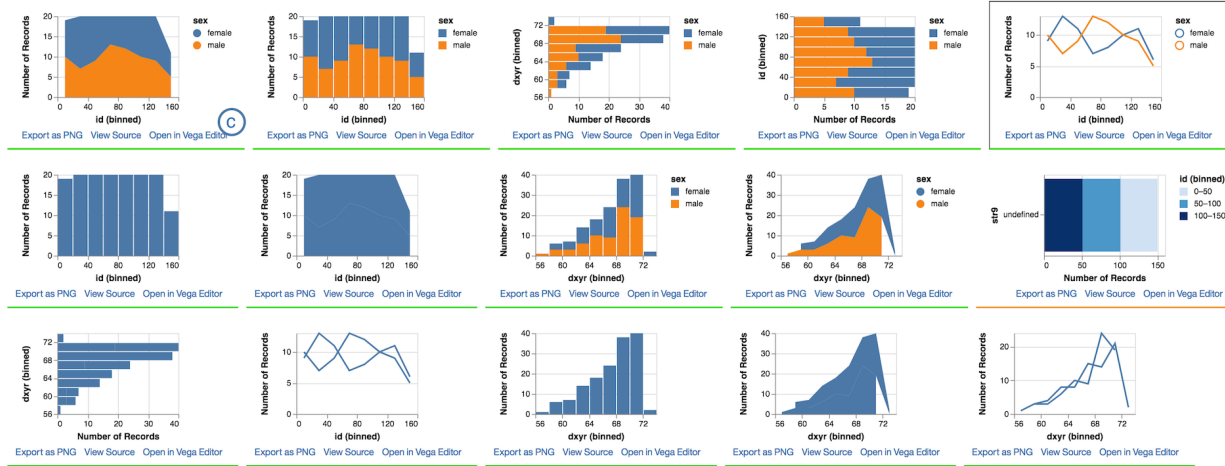
Figure 6: Data2Vis qualitative evaluation interface with results from beam search. (a) A user can load a random dataset from the RDdataset collection or paste a dataset (JSON format) and select "Generate." (b) User can paste a JSON dataset and select "Generate"(c) Data2Vis generates Vega-Lite specifications using beam search (beam width = 15 in this case) based on the dataset. The user can modify and iterate on any of the visualizations using the Vega-Lite editor. Highlights below each visualization represents cases of valid specifications ▪ and incomplete specifications ▪ where the model attempts to use variables not in the dataset (phantom variables).


Figure 7: Examples of visualizations where the model has learned common selection patterns and leverages concepts such as responses (yes, no) and sex (male, female).

their needs. This is especially true when the structure of the data being analyzed is unknown or unusual and effective exploration of the data requires complex transforms as well as deep domain expertise. However, Data2Vis can contribute to this process by "jumpstarting" the visualization process—first by generating a set of valid visualization specifications and *seeding* the creativity process with these initial visualizations. Analysts can initialize their visualization tasks with Data2Vis and iteratively *correct* its content while generating intermediate visualizations.

7.2 Limitations

Field Selection and Transformation The current version of our model has limitations which occur in about 15-20% of tests. First, the model occasionally selects what we refer to as a phantom field (a field that does not exist in the input dataset) as part of the visualization spec it generates (Figure 6). While plots are still realized in some cases despite this error (Vega-Lite incorporates good defaults), the affected axis is not interpretable. Another limitation of the model is observed in selecting fields (attributes) of the input data to visualize — the model sometime selects fields that are unintuitive or have little information value. For example, a frequency plot of grouped longitude and latitude values does not provide much infor-

mation. Finally, the model generates relatively simple visualizations — univariate plots (which can serve as data field summaries) and bivariate plots. It is unable to apply complex transforms, use multiple variables.

Training Data While further experimentation is required, our intuition is that the limitations mentioned above reflect limitations in both the size and diversity of our training data. Our goal with Data2Vis was to evaluate the viability of machine translation in generating valid visualization specifications, we have conducted our experiments with a relatively small dataset (4300 examples sampled to 215,000 training pairs). While our current results provide insights, we believe a larger and more diversified training dataset will improve learning and model generalization. Another limitation with our training data is related to our training pair generation strategy. Currently, we construct our source tokens from a single row from a dataset which is then preprocessed before training. While this approach shortens the input sequence length, a requirement for us to efficiently train our model, the model can only learn properties of each field (e.g. length, content, numeric type, string type) as opposed to properties of the distribution of the field (e.g mean, range, categories etc.) which encode useful signals for data visualization.

Figure 8: Examples where the model has learned to generate univariate plots that summarize fields selected from the dataset.

Figure 9: A comparison of visualizations generated by Data2vis(a) and Voyager 2(b) given the same race dataset.

7.3 Future Work

Eliciting More Training Data Naturally, addressing limitations with our training data constitutes the next step for future work. We plan to conduct a structured data collection aimed at generating visualization examples across a large number of datasets, visualization types (bar, point, area, chart etc), transforms, complexity (number of variables), interactions and visualization languages. We will also explore strategies to improve the training process that guide the model towards learning properties of the distribution for a given field.

Extending Data2Vis to Generate Multiple Plausible Visualizations Data2Vis is currently implemented as a sequence to sequence translation model. Sequence models work very well for domains where it is desirable to have fixed mappings of input sequences to output sequences (text summarization, image captioning, language translation, etc). It is generally expected that a sentence in one language always maps to the same sentence in a different language, and acceptable if a passage always maps to the same summary or an image to the same caption. However, when applied to the task

of data visualization, it is desirable that input data maps to *multiple* valid visualizations. In the current work, we address this by exploiting beam search decoding to generate multiple visualizations based on a single dataset. A related avenue for future work is to explore generative models that can learn a probability distribution of effective visualizations, enabling *one to many sequence* mappings between data and visualization specifications through sampling.

Targeting Additional Grammars Building on results from Data2Vis, important next steps also include efforts to train models that can map input data to multiple different visualization specification languages, including ggplot2, given a dataset. This line of research may also explore training models that learn direct mappings between different visualization specification languages, enabling visualization specification reuse across languages and platforms.

Natural Language and Visualization Specification We propose the exploration of models that generate visualizations conditioned on natural language text in addition to datasets. A potential approach is to first explore how users might describe or express visualizations for a given dataset and use this knowledge in generation of triplets—

natural language description, data sequence, and visualization specification. These data points can then be leveraged in training a model that learns to generate visualizations based on natural language descriptions. These models can extend the expressive capabilities of existing systems that integrate multimodal interactions and visualizations for exploring data. Conversely, we can use textual descriptions of visualizations to automatically generate captions for them, akin to image caption generation (e.g., [36, 70, 80]).

8 CONCLUSION

The history of data visualization is rich with work that treats visualization from a linguistic perspective. Bertin systematized data visualization as "a language for the eye" [7]. Adopting Bertin's analogy, Mackinlay [47] viewed visualizations as sentences of a graphical language and formalized a model based on "expressiveness" and "effectiveness" criteria, borrowing concepts from formal languages. Subsequent research also introduced various "grammars" of visualization specification.

We significantly extend this earlier perspective and formulate data visualization as a sequence to sequence translation problem where we translate data specifications to visualization specifications. We train a deep sequence to sequence model and demonstrate its efficacy generating univariate and bivariate plots. We also identify initial failure conditions, offer ideas for their remediation and an agenda for future work.

It is our belief that the problem formulation and model presented in this work represents an appropriate baseline for future work in automated generation of visualizations using deep learning approaches. Our approach sets the stage for systems that learn to generate visualizations at scale with implications for the development of guided visual data exploration systems.

9 ACKNOWLEDGMENTS

REFERENCES

[1] M. M. Alan. *How Maps Work: Representation, Visualization, and Design*. Guilford Press, 1995.

[2] D. Asimov. The grand tour: A tool for viewing multidimensional data. *SIAM J. Sci. Stat. Comput.*, 6(1), 1985.

[3] D. Bahdanau, K. Cho, and Y. Bengio. Neural Machine Translation by Jointly Learning to Align and Translate. *ICLR 2015*, sep 2014.

[4] M. Balog, A. L. Gaunt, M. Brockschmidt, S. Nowozin, and D. Tarlow. Deepcoder: Learning to write programs. *CoRR*, abs/1611.01989, 2016.

[5] S. B. Barbara. Stimulus processing models from psychology: can we use them in cartography? *The American Cartographer*, 9:155–167, 1982.

[6] T. Beltramelli. pix2code: Generating code from a graphical user interface screenshot. *arXiv preprint arXiv:1705.07962*, 2017.

[7] J. Bertin. *Semiology of Graphics*. University of Wisconsin Press, 1983.

[8] P. Bojanowski, A. Joulin, and T. Mikolov. Alternative structures for character-level rnns. *CoRR*, abs/1511.06303, 2015.

[9] M. Bostock and J. Heer. Protovis: A graphical toolkit for visualization. *IEEE TVCG (Proc. InfoVis)*, 2009.

[10] M. Bostock, V. Ogievetsky, and J. Heer. D3: Data-driven documents. *IEEE TVCG (Proc. InfoVis)*, 2011.

[11] F. Bouali, A. Guettala, and G. Venturini. VizAssist: an interactive user assistant for visual data mining. *Vis. Comput.*, 32(11):1447–1463, 2016.

[12] N. Boulanger-Lewandowski, Y. Bengio, and P. Vincent. Modeling temporal dependencies in high-dimensional sequences: Application to polyphonic music generation and transcription. *arXiv:1206.6392*, 2012.

[13] D. Britz, A. Goldie, T. Luong, and Q. Le. Massive Exploration of Neural Machine Translation Architectures. *ArXiv e-prints*, Mar. 2017.

[14] S. M. Casner. Task-analytic approach to the automated design of graphic presentations. *ACM Trans. Graphics*, 10(2):111–151, 1991.

[15] X. Chen, C. Liu, and D. Song. Tree-to-tree neural networks for program translation. *CoRR*, abs/1802.03691, 2018.

[16] K. Cho, B. van Merrienboer, Ç. Gülçehre, F. Bougares, H. Schwenk, and Y. Bengio. Learning phrase representations using RNN encoder-decoder for statistical machine translation. *CoRR*, abs/1406.1078, 2014.

[17] W. S. Cleveland and R. McGill. Graphical perception: Theory, experimentation, and application to the development of graphical methods. *J. Amer. Statist. Assoc.*, 79(387):531–554, 1984.

[18] Ç. Demiralp, P. J. Haas, S. Parthasarathy, and T. Pedapati. Foresight: Recommending visual insights. *Proc. VLDB Endow.*, 10(12):1937–1940, 2017.

[19] Ç. Demiralp, C. Scheidegger, G. Kindlmann, D. Laidlaw, and J. Heer. Visual embedding: A model for visualization. *IEEE CG&A*, 2014.

[20] J. Devlin, J. Uesato, S. Bhupatiraju, R. Singh, A. Mohamed, and P. Kohli. Robustfill: Neural program learning under noisy I/O. *CoRR*, abs/1703.07469, 2017.

[21] V. Dibia. Data2vis: Automatic generation of data visualizations using sequence to sequence recurrent neural networks. `https://github.com/victordibia/data2vis`, 2017.

[22] V. Dibia and Ç. Demiralp. Data2Vis: Automatic generation of visualizations using sequence-to-sequence recurrent neural networks, 2018.

[23] L. Dong and M. Lapata. Language to Logical Form with Neural Attention. jan 2016.

[24] D. Eck and J. Schmidhuber. A first look at music composition using lstm recurrent neural networks. *Istituto Dalle Molle Di Studi Sull Intelligenza Artificiale*, 103, 2002.

[25] M. A. Fisherkeller, J. H. Friedman, and J. W. Tukey. Prim-9: An interactive multidimensional data display and analysis system. In *Proc. Fourth International Congress for Stereology*, 1974.

[26] F. A. Gers, J. Schmidhuber, and F. Cummins. Learning to forget: Continual prediction with lstm. *Neural Computation*, 12:2451–2471, 1999.

[27] D. Gotz and Z. Wen. Behavior-driven visualization recommendation. In *ACM IUI*, pp. 315–324, 2009.

[28] A. Graves. Sequence Transduction with Recurrent Neural Networks, nov 2012.

[29] D. Ha and D. Eck. A neural representation of sketch drawings. *CoRR*, abs/1704.03477, 2017.

[30] P. Hanrahan. Vizql: a language for query, analysis and visualization. In *Proceedings of the 2006 ACM SIGMOD international conference on Management of data*, pp. 721–721. ACM, 2006.

[31] S. Hochreiter and J. Schmidhuber. Long short-term memory. 9:1735–80, 12 1997.

[32] K. Z. Hu, M. A. Bakker, S. Li, T. Kraska, and C. A. Hidalgo. Vizml: A machine learning approach to visualization recommendation, 2018.

[33] S. Jean, K. Cho, R. Memisevic, and Y. Bengio. On using very large target vocabulary for neural machine translation. *arXiv:1412.2007*, 2014.

[34] D. D. Johnson. Generating polyphonic music using tied parallel networks. In *International Conference on Evolutionary and Biologically Inspired Music and Art*, pp. 128–143. Springer, 2017.

[35] N. Kalchbrenner and P. Blunsom. Recurrent continuous translation models. In *Proceedings of the 2013 Conference on Empirical Methods in Natural Language Processing*, pp. 1700–1709, 2013.

[36] A. Karpathy and L. Fei-Fei. Deep visual-semantic alignments for generating image descriptions. In *Proc. IEEE CVPR*, pp. 3128–3137, 2015.

[37] A. Karpathy, J. Johnson, and F. Li. Visualizing and understanding recurrent networks. *CoRR*, abs/1506.02078, 2015.

[38] T. Karras, T. Aila, S. Laine, and J. Lehtinen. Progressive growing of gans for improved quality, stability, and variation. *CoRR*, abs/1710.10196, 2017.

[39] A. Key, B. Howe, D. Perry, and C. Aragon. VizDeck. In *ACM SIGMOD*, vol. 681–684, 2012.

[40] G. Kindlmann and C. Scheidegger. An algebraic process for visualization design. *IEEE TVCG*, 20:2181–2190, 2014.

[41] S. Lewandowsky and I. Spence. Discriminating strata in scatterplots. *Journal of American Statistical Association*, 84(407):682–688, 1989.

[42] W. Ling, E. Grefenstette, K. M. Hermann, T. Kociský, A. Senior, F. Wang, and P. Blunsom. Latent predictor networks for code generation. *CoRR*, abs/1603.06744, 2016.

[43] Y. Luo, X. Qin, N. Tang, G. Li, and X. Wang. Deepeye: Creating good data visualizations by keyword search. In *Proceedings of the 2018 International Conference on Management of Data*, SIGMOD, pp. 1733–1736, 2018.

[44] M.-T. Luong and C. D. Manning. Achieving Open Vocabulary Neural Machine Translation with Hybrid Word-Character Models. *ACL 2016*, apr 2016.

[45] M.-T. Luong, H. Pham, and C. D. Manning. Effective Approaches to Attention-based Neural Machine Translation. *EMNLP 2015*, aug 2015.

[46] M.-T. Luong, I. Sutskever, Q. V. Le, O. Vinyals, and W. Zaremba. Addressing the rare word problem in neural machine translation. *arXiv:1410.8206*, 2014.

[47] J. Mackinlay. Automating the design of graphical presentations of relational information. *ACM Trans. Graphics*, 5(2):110–141, 1986.

[48] J. Mackinlay, P. Hanrahan, and C. Stolte. Show me: Automatic presentation for visual analysis. *IEEE TVCG*, 13(6):1137–1144, 2007.

[49] D. Moritz, C. Wang, G. Nelson, H. Lin, A. M. Smith, B. Howe, and J. Heer. Formalizing visualization design knowledge as constraints: Actionable and extensible models in draco. *IEEE TVCG (Proc. InfoVis)*, 2019.

[50] R. Nallapati, B. Zhou, C. N. dos Santos, C. Gulcehre, and B. Xiang. Abstractive Text Summarization Using Sequence-to-Sequence RNNs and Beyond. *The SIGNLL Conference on Computational Natural Language Learning (CoNLL), 2016*, feb 2016.

[51] E. Parisotto, A. Mohamed, R. Singh, L. Li, D. Zhou, and P. Kohli. Neuro-symbolic program synthesis. *CoRR*, abs/1611.01855, 2016.

[52] J. Poco and J. Heer. Reverse-engineering visualizations: Recovering visual encodings from chart images. *Computer Graphics Forum (Proc. EuroVis)*, 2017.

[53] S. Reed, Z. Akata, X. Yan, L. Logeswaran, B. Schiele, and H. Lee. Generative adversarial text to image synthesis. *arXiv preprint arXiv:1605.05396*, 2016.

[54] S. F. Roth, J. Kolojechick, J. Mattis, and J. Goldstein. Interactive graphic design using automatic presentation knowledge. In *ACM Human Factors in Computing Systems (CHI)*, 1994.

[55] B. Saket, A. Endert, and Ç. Demiralp. Task-based effectiveness of basic visualizations. *IEEE TVCG*, 2018.

[56] B. Saket, D. Moritz, H. Lin, V. Dibia, Ç. Demiralp, and J. Heer. Beyond heuristics: Learning visualization design, 2018.

[57] A. Satyanarayan, D. Moritz, K. Wongsuphasawat, and J. Heer. Vega-lite: A grammar of interactive graphics. *IEEE TVCG (Proc. InfoVis)*, 2017.

[58] A. Satyanarayan, R. Russell, J. Hoffswell, and J. Heer. Reactive vega: A streaming dataflow architecture for declarative interactive visualization. *IEEE TVCG (Proc. InfoVis)*, 2016.

[59] M. Schuster and K. K. Paliwal. Bidirectional recurrent neural networks. *IEEE Transactions on Signal Processing*, 45(11):2673–2681, 1997. doi: 10.1109/78.650093

[60] R. Sennrich, B. Haddow, and A. Birch. Edinburgh neural machine translation systems for wmt 16. *arXiv preprint arXiv:1606.02891*, 2016.

[61] J. Seo and B. Shneiderman. A rank-by-feature framework for unsupervised multidimensional data exploration using low dimensional projections. In *Procs. InfoVis*, pp. 65–72, 2004.

[62] R. N. Shepard. Toward a universal law of generalization for psychological science. *Science*, 237:1317–1323, 1987.

[63] T. Siddiqui, A. Kim, J. Lee, K. Karahalios, and A. Parameswaran. Effortless data exploration with Zenvisage. *PVLDB*, 10(4):457–468, 2016.

[64] C. Stolte, D. Tang, and P. Hanrahan. Polaris: a system for query, analysis, and visualization of multidimensional relational databases. *IEEE TVCG*, 8(1):52–65, 2002. doi: 10.1109/2945.981851

[65] Sumit Chopra, Michael Auli, and Alexander M. Rush. Abstractive sentence summarization with attentive recurrent neural networks. In *Proceedings of NAACL-HLT*, 2016.

[66] I. Sutskever, O. Vinyals, and Q. V. Le. Sequence to sequence learning with neural networks. *CoRR*, abs/1409.3215, 2014.

[67] B. Tang, S. Han, M. L. Yiu, R. Ding, and D. Zhang. Extracting top-k insights from multi-dimensional data. In *ACM SIGMOD*, pp. 1509–1524, 2017.

[68] M. Vartak, S. Huang, T. Siddiqui, S. Madden, and A. Parameswaran. Towards visualization recommendation systems. In *DSIA Workshop*, 2015.

[69] M. Vartak, S. Rahman, S. Madden, A. Parameswaran, and N. Polyzotis. SeeDB: Efficient data-driven visualization recommendations to support visual analytics. *PVLDB*, 8(13):2182–2193, 2015.

[70] O. Vinyals, A. Toshev, S. Bengio, and D. Erhan. Show and tell: A neural image caption generator. In *Computer Vision and Pattern Recognition (CVPR), 2015 IEEE Conference on*, pp. 3156–3164. IEEE, 2015.

[71] H. Wickham. A layered grammar of graphics. *Journal of Computational and Graphical Statistics*, 19(1):3–28, 2010.

[72] L. Wilkinson. *The Grammar of Graphics*. Springer, 1st ed., 1999.

[73] L. Wilkinson, A. Anand, and R. Grossman. Graph-theoretic scagnostics. In *Proc. InfoVis*, pp. 157–164, 2005.

[74] G. Wills. Brunel v2.5. https://github.com/Brunel-Visualization/Brunel, 2017. Accessed: 2018-04-04.

[75] G. Wills and L. Wilkinson. AutoVis: Automatic visualization. *Info. Visual.*, 9(1):47–69, 2008.

[76] K. Wongsuphasawat, D. Moritz, A. Anand, J. Mackinlay, B. Howe, and J. Heer. Towards a general-purpose query language for visualization recommendation. In *Procs. HILDA*, 2016.

[77] K. Wongsuphasawat, D. Moritz, A. Anand, J. Mackinlay, B. Howe, and J. Heer. Voyager: Exploratory analysis via faceted browsing of visualization recommendations. *IEEE TVCG*, 22(1):649–658, 2016.

[78] K. Wongsuphasawat, Z. Qu, D. Moritz, R. Chang, F. Ouk, A. Anand, J. Mackinlay, B. Howe, and J. Heer. Voyager 2: Augmenting visual analysis with partial view specifications. In *ACM CHI*, 2017.

[79] Y. Wu, M. Schuster, Z. Chen, Q. V. Le, M. Norouzi, W. Macherey, M. Krikun, Y. Cao, Q. Gao, K. Macherey, et al. Google's neural machine translation system: Bridging the gap between human and machine translation. *arXiv:1609.08144*, 2016.

[80] K. Xu, J. Ba, R. Kiros, K. Cho, A. Courville, R. Salakhudinov, R. Zemel, and Y. Bengio. Show, attend and tell: Neural image caption generation with visual attention. In *Proc. ICML*, pp. 2048–2057, 2015.

[81] P. Yin and G. Neubig. A syntactic neural model for general-purpose code generation. *CoRR*, abs/1704.01696, 2017.

[82] V. Zhong, C. Xiong, and R. Socher. Seq2SQL: Generating Structured Queries from Natural Language using Reinforcement Learning. aug 2017.

Figure 10: Examples of visualizations generated with beam search.